

# Viper Mortality

Vivi B. Ngo

```
# Packages you definitely need
library(ggplot2)
library(dplyr)
library(ranger) # for looking up help on Conceptual Problem
library(MASS) # see Applied Problem 2
library(parsnip)
library(rsample)
# Please load any other packages you need either all in this chunk
# or in the chunk you need it for
library(keras)
library(tensorflow)
library(ISLR2)
library(esquisse)

# Data you need for applied problems
SoCalRent <- readr::read_csv("SoCalRent.csv")
viper_train <- readr::read_csv("viper_train.csv")
viper_test <- readr::read_csv("viper_test.csv")
```

## Snakebite Mortality

Gopalakrishnan and colleagues (2022) attempted to build a model to predict mortality from snakebite based on predictors that could be measured within 48 hours of admission to a hospital or clinic.

The `viper_train` dataset on Canvas contains information about 239 patients from southern India who were diagnosed with having been poisoned by a viper bite. These patients were used to fit their model. The `viper_test` dataset on Canvas contains the same information about a separate set of 140 patients, who were used to validate the model. Please see the `viper_dictionary` file on Canvas for an explanation of each of the variables.

- Changing the `Outcome` variable appropriately, or creating a new response based on `Outcome`, to reflect that this is a classification problem

```
viper_test$Outcome <- factor(viper_test$Outcome)
viper_train$Outcome <- factor(viper_train$Outcome)

#viper_test$Sex <- factor(viper_test$Sex)
#viper_train$Sex <- factor(viper_train$Sex)
#viper_test$Outcome<-ifelse(viper_test$Outcome=="0",0,1)
```

```
#viper_train$Outcome<-ifelse(viper_train$Outcome=="0",0,1)
```


```
viper_test$Sex<-ifelse(viper_test$Sex=="M",1,0)  
viper_train$Sex<-ifelse(viper_train$Sex=="M",1,0)
```

- Performing exploratory data analysis on the training set and documenting your findings

```
#eda on response of interest: outcome  
viper_train %>% group_by(Outcome)%>%count()
```

```
## # A tibble: 2 x 2  
## # Groups:   Outcome [2]  
##   Outcome     n  
##   <fct>   <int>  
## 1 0       187  
## 2 1        52
```


```
#we have 187 people who lived a snake bite and 52 who died in this training set  
ggplot(viper_train, aes(x=Outcome)) + geom_bar()
```



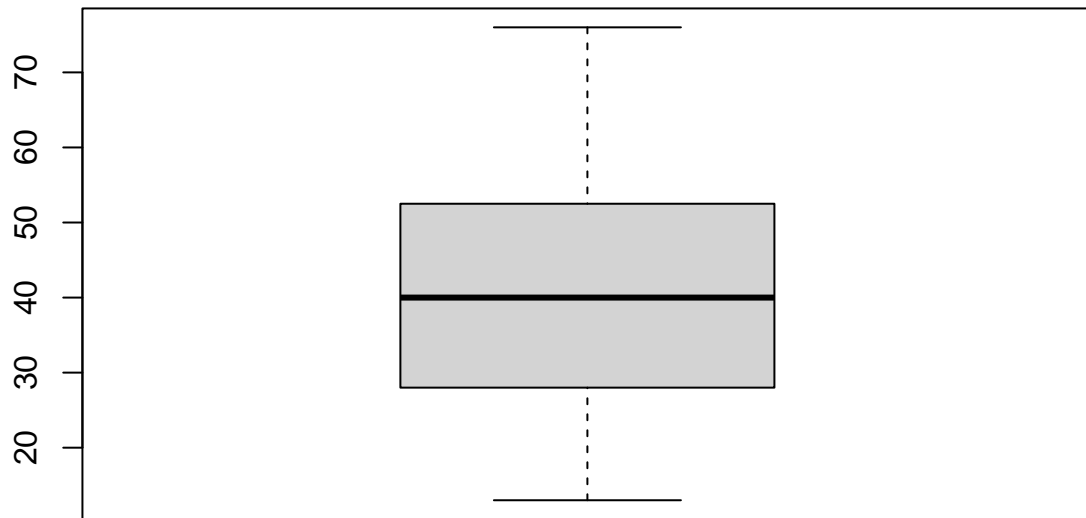
```
#eda on predictors  
#77 females and 162 males in the training set  
viper_train %>% group_by(Sex)%>%count()
```

```
## # A tibble: 2 x 2
## # Groups:   Sex [2]
##   Sex     n
##   <dbl> <int>
## 1     0    77
## 2     1   162
```

```
ggplot(viper_train, aes(x=Sex)) + geom_bar()
```

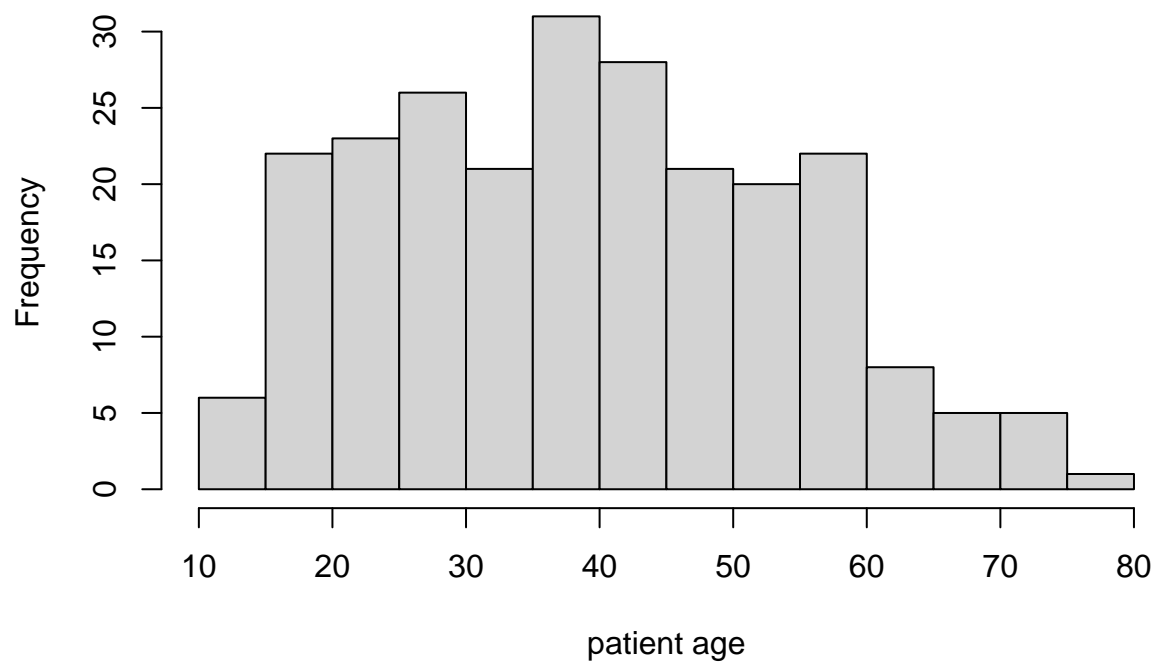


```
#age
age_boxplot <- boxplot(viper_train$Age)
```



```
age_hist <- hist(viper_train$Age, xlab = 'patient age')
```

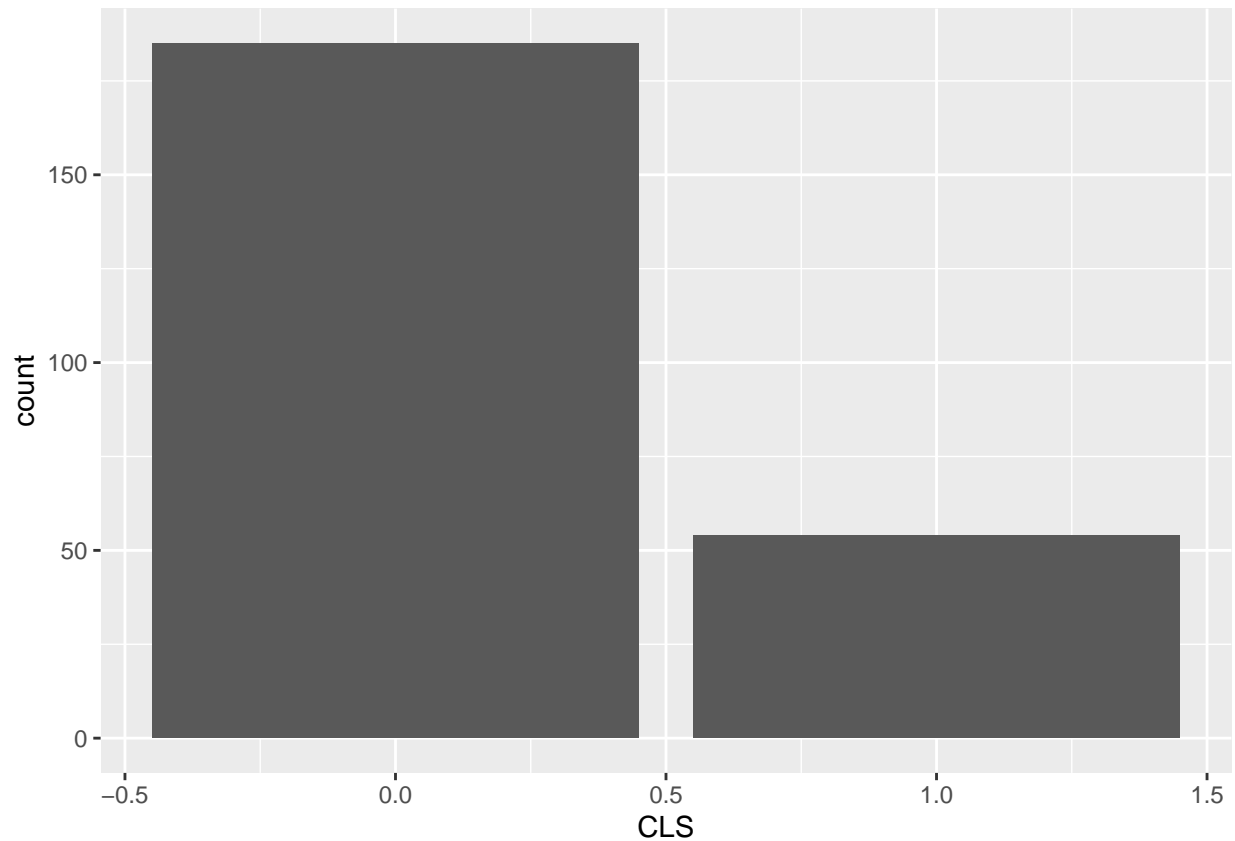
**Histogram of viper\_train\$Age**



```
#CLS capillary leak syndrome  
# no(0) = 185, yes(1) = 54  
viper_train %>% group_by(CLS)%>%count()
```

```
## # A tibble: 2 x 2  
## # Groups:   CLS [2]  
##   CLS     n  
##   <dbl> <int>  
## 1     0  185  
## 2     1   54
```


```
ggplot(viper_train, aes(x=CLS)) + geom_bar()
```



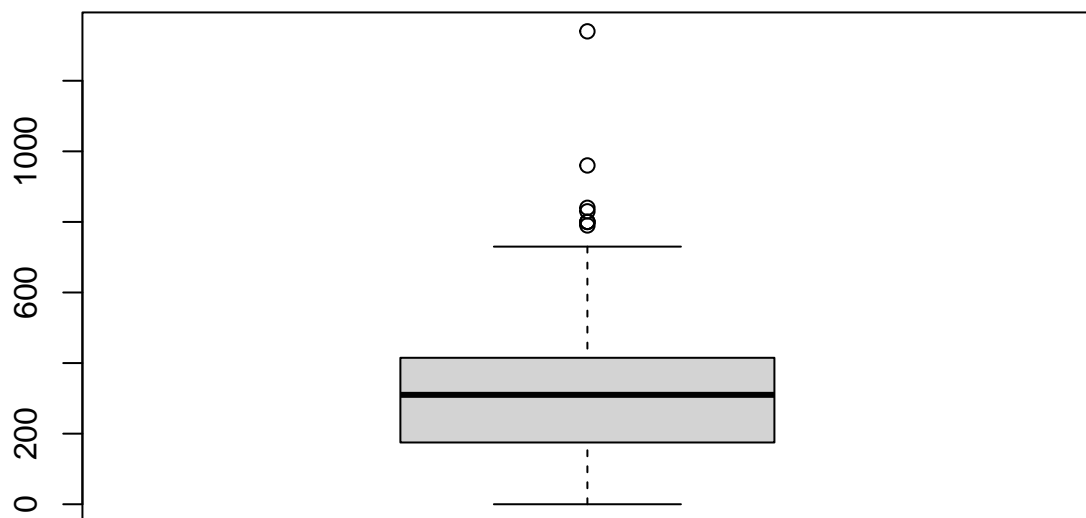
```
#Bleeding  
# no bleeding = 107, yes bleeding = 132  
viper_train %>% group_by(Bleeding)%>%count()
```

```
## # A tibble: 2 x 2  
## # Groups:   Bleeding [2]  
##   Bleeding     n  
##   <dbl> <int>  
## 1     0   107  
## 2     1   132
```

```
ggplot(viper_train, aes(x=Bleeding)) + geom_bar()
```




```
#ASVTotal  
asv_boxplot <- boxplot(viper_train$ASVTotal)
```




```
asv_hist <- hist(viper_train$ASVTotal, xlab = 'Total amount of antivenom administered')
```



**Histogram of viper\_train\$ASVTotal**




```
#platelets  
plate_boxplot <- boxplot(viper_train$Platelets)
```



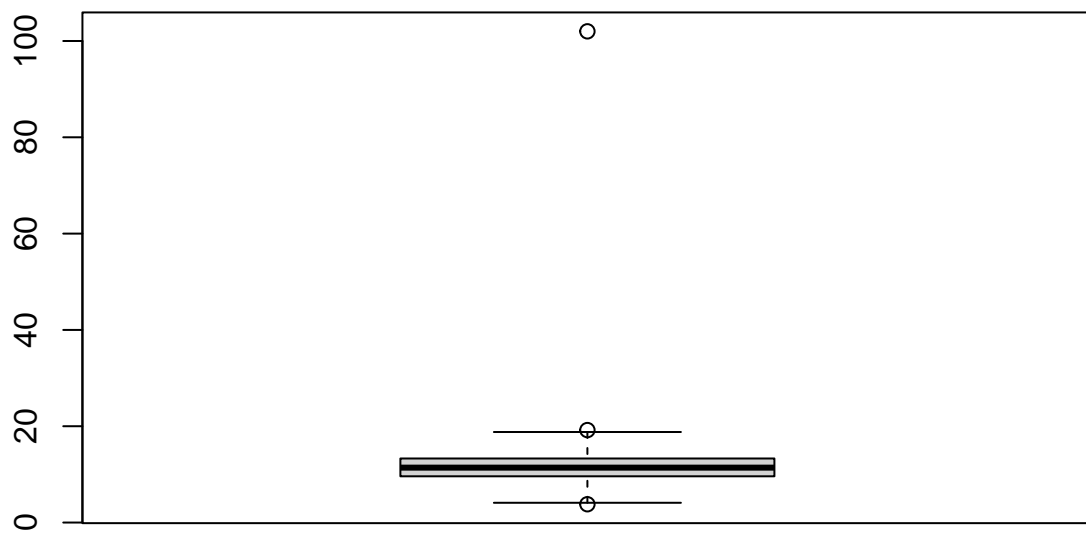
```
plate_hist <- hist(viper_train$Platelets, xlab = 'Platelet count at  
admission')
```

**Histogram of viper\_train\$Platelets**



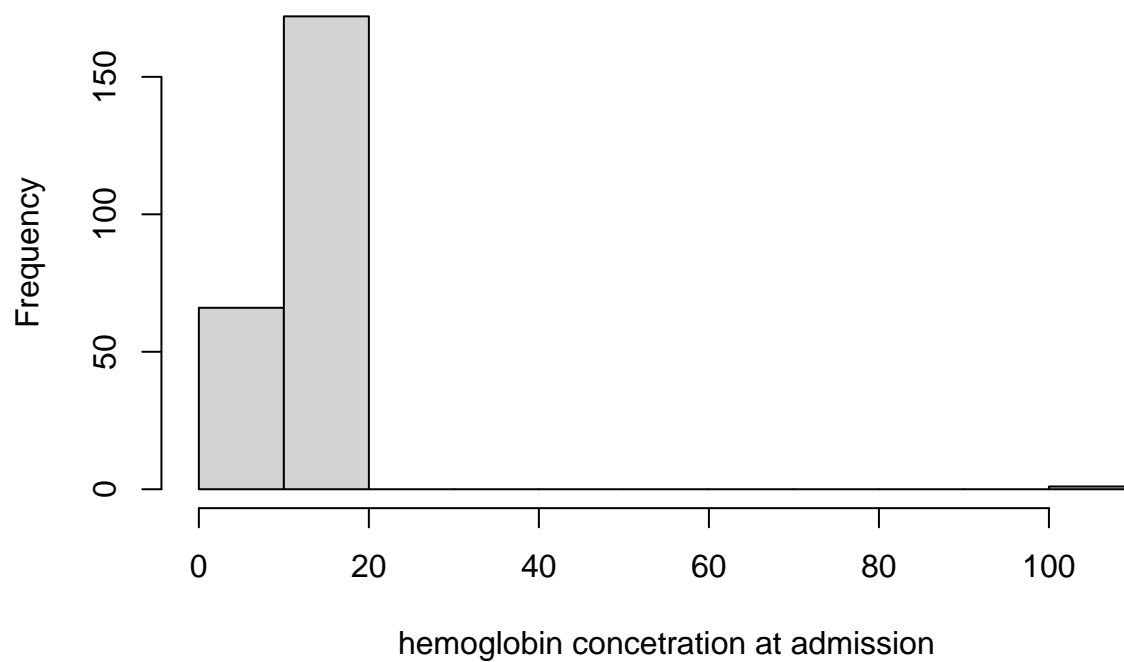
*#Hb*

```
hb_boxplot <- boxplot(viper_train$Hb)
```

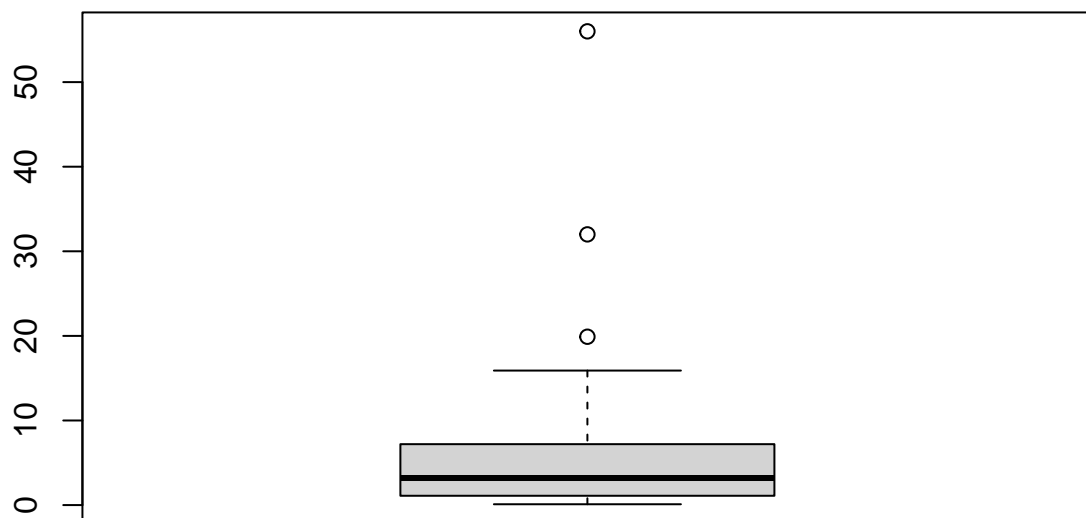


```
hb_hist <- hist(viper_train$Hb, xlab = 'hemoglobin concetration at admission')
```

**Histogram of viper\_train\$Hb**

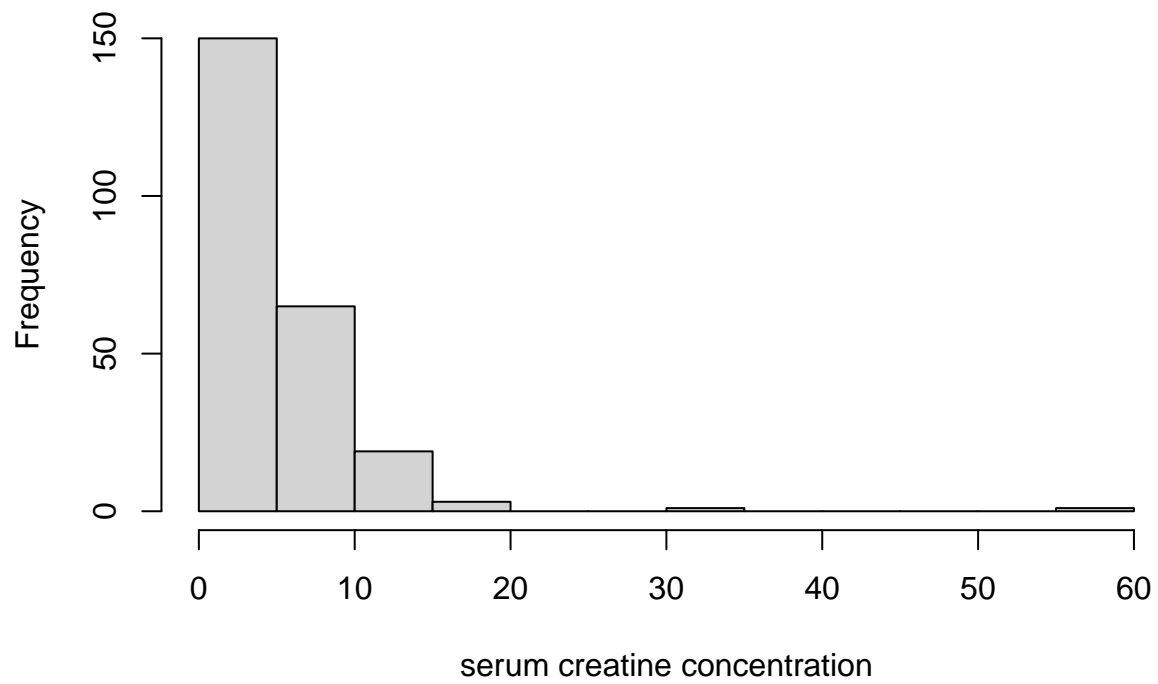


```
#creatinine  
creatinine_boxplot <- boxplot(viper_train$Creatinine)
```

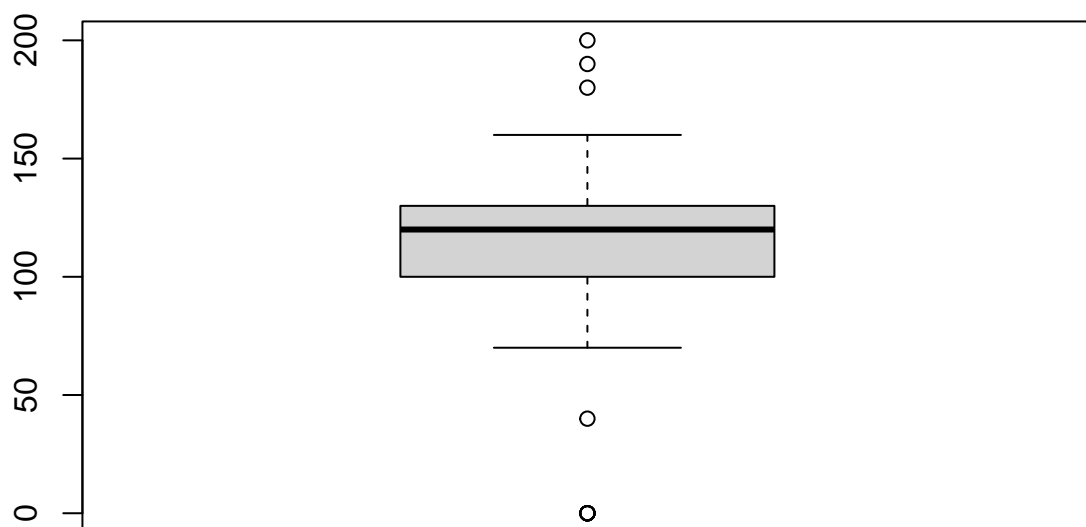


```
creatinine_hist <- hist(viper_train$Creatine, xlab = 'serum creatine concentration')
```

**Histogram of viper\_train\$Creatine**



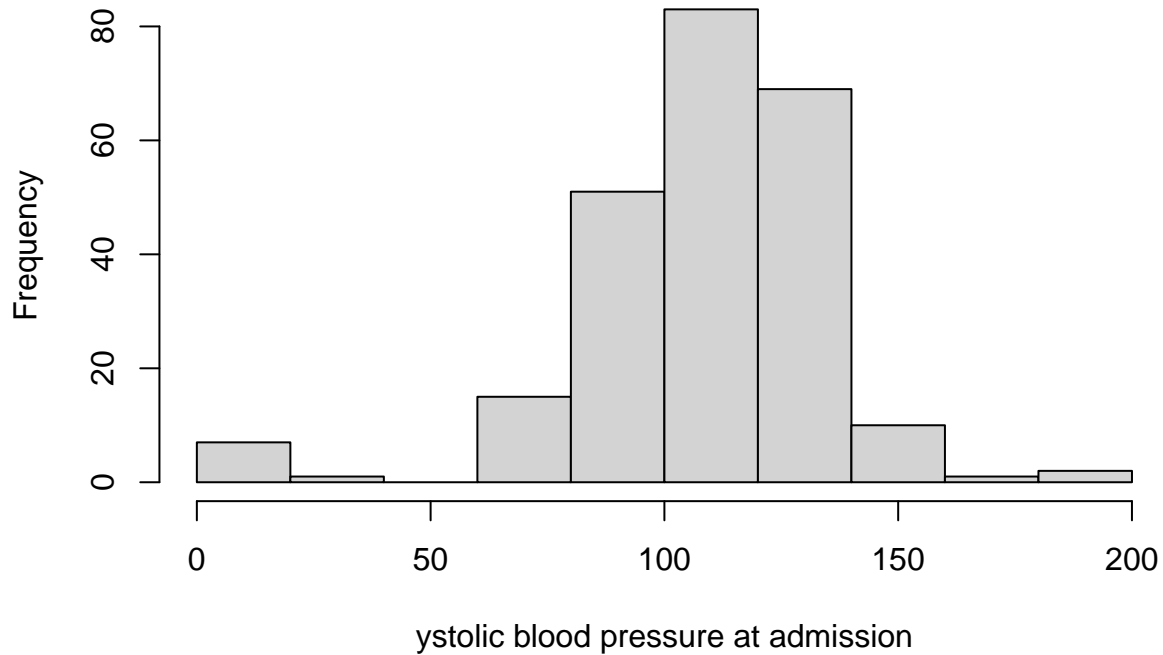
```
#blood pressure  
#is it possible to have 200bp?  
#can dead people have more than 0 bp?  
bp_boxplot <- boxplot(viper_train$BloodPressure)
```



```
bp_hist <- hist(viper_train$BloodPressure, xlab = 'ystolic blood pressure at admission')
```



## Histogram of viper\_train\$BloodPressure



- Performing backward stepwise selection using BIC as a selection criterion and obtaining an estimate of the test error rate on the selected logistic regression model (HINT: `regsubsets` will not work for logistic regression, but you can use the `stepAIC` function in the `MASS` package instead to do stepwise selection. The output of `stepAIC` will be a `glm` object, which means you can do prediction without the workarounds)

```
library(MASS)
library(tidyverse)
```

```
## -- Attaching packages ----- tidyverse 1.3.1 --
```

```
## v tibble  3.1.6      v purrr   0.3.4
## v tidyr   1.1.4      v stringr 1.4.0
## v readr   2.1.2      v forcats 0.5.1
```

```
## -- Conflicts ----- tidyverse_conflicts() --
```

```
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()
## x MASS::select()  masks dplyr::select()
```

```
set.seed(10)
```

```
outcome_test <- viper_test$Outcome
```

```
bglm <- glm(Outcome~.,data=viper_train, family = 'binomial')
backward <- stepAIC(bglm, data=viper_train,method='backward')
```

```
## Start:  AIC=114.95
## Outcome ~ Age + Sex + CLS + Bleeding + ASVTotal + Platelets +
##      Hb + Creatine + BloodPressure
##
##              Df Deviance    AIC
## - Creatine      1   95.008 113.01
## - Age            1   95.263 113.26
## - Platelets      1   95.530 113.53
## - Sex            1   96.910 114.91
## <none>           94.951 114.95
## - Hb             1   98.794 116.79
## - ASVTotal       1  102.381 120.38
## - BloodPressure  1  117.582 135.58
## - CLS            1  119.912 137.91
## - Bleeding       1  122.572 140.57
##
## Step:  AIC=113.01
## Outcome ~ Age + Sex + CLS + Bleeding + ASVTotal + Platelets +
##      Hb + BloodPressure
##
##              Df Deviance    AIC
## - Age            1   95.339 111.34
## - Platelets      1   95.904 111.90
## - Sex            1   96.982 112.98
## <none>           95.008 113.01
## - Hb             1   98.933 114.93
## - ASVTotal       1  102.650 118.65
## - BloodPressure  1  118.002 134.00
## - CLS            1  120.162 136.16
## - Bleeding       1  122.963 138.96
##
## Step:  AIC=111.34
## Outcome ~ Sex + CLS + Bleeding + ASVTotal + Platelets + Hb +
##      BloodPressure
##
##              Df Deviance    AIC
## - Platelets      1   96.233 110.23
## <none>           95.339 111.34
## - Sex            1   97.724 111.72
## - Hb             1   99.717 113.72
## - ASVTotal       1  102.727 116.73
## - BloodPressure  1  118.310 132.31
## - CLS            1  120.870 134.87
## - Bleeding       1  124.021 138.02
##
## Step:  AIC=110.23
## Outcome ~ Sex + CLS + Bleeding + ASVTotal + Hb + BloodPressure
##
##              Df Deviance    AIC
## <none>           96.233 110.23
```

```
## - Sex          1   98.405 110.41
## - Hb           1  101.192 113.19
## - ASVTotal     1  103.579 115.58
## - BloodPressure 1  121.528 133.53
## - CLS          1  125.266 137.27
## - Bleeding     1  127.859 139.86
```

```
back_pred <- predict(backward, viper_test)
```

```
#back_error <- sum((outcome_test/sum(backward$residuals))^2)
```

- Fitting at least two of the following four types of models and obtaining an estimate of the test error rate on each model: k-nearest neighbors, generative models (LDA/QDA/naive Bayes), tree-based methods (bagging/random forests/boosting), neural networks

```
set.seed(10)
```

```
#viper_test$Sex<-ifelse(viper_test$Sex=="M",1,0)
#viper_train$Sex<-ifelse(viper_train$Sex=="M",1,0)
#boosted tree
```

```
x_scale_test <- data.frame(
  outcome = as.numeric(viper_test$Outcome),
  sex = as.numeric(viper_test$Sex),
  age = scale(viper_test$Age),
  cls = scale(viper_test$CLS),
  bleeding = scale(viper_test$Bleeding),
  asvTotal = scale(viper_test$ASVTotal),
  platelets = scale(viper_test$Platelets),
  hb = scale(viper_test$Hb),
  creatine = scale(viper_test$Creatine),
  bp = scale(viper_test$BloodPressure)
) %>% as.matrix()
```

```
x_scale_train <- data.frame(
  outcome = as.numeric(viper_train$Outcome),
  sex = as.numeric(viper_train$Sex),
  age = scale(viper_train$Age),
  cls = scale(viper_train$CLS),
  bleeding = scale(viper_train$Bleeding),
  asvTotal = scale(viper_train$ASVTotal),
  platelets = scale(viper_train$Platelets),
  hb = scale(viper_train$Hb),
  creatine = scale(viper_train$Creatine),
  bp = scale(viper_train$BloodPressure)
) %>% as.matrix()
```

```

#viper_test$Sex <- as.numeric(viper_test$Sex)

#1/2 (fast n dirty random forest)
library(ranger)
viper_rf<- ranger(Outcome~., data=viper_train, importance = "permutation", seed=758)
viper_rf

## Ranger result
##
## Call:
## ranger(Outcome ~ ., data = viper_train, importance = "permutation",      seed = 758)
##
## Type:                                Classification
## Number of trees:                      500
## Sample size:                          239
## Number of independent variables:      9
## Mtry:                                  3
## Target node size:                     1
## Variable importance mode:             permutation
## Splitrule:                             gini
## OOB prediction error:                 10.88 %

#estimate the test MSE :3333333
#(find the default test)
viper_class <- predict(viper_rf, data=viper_test)$predictions

#predict each tree and then avg the preds
viper_predictions <- predict(viper_rf, data = viper_test, predict.all = TRUE)

rf_probs <- apply(viper_predictions$predictions-1,1,mean)

#neural net
library(keras)

nn_1layer <- keras_model_sequential() %>%
  layer_dense(units = 10, activation = "relu",
  input_shape = ncol(x_scale_train)) %>%
  layer_dropout(rate = 0.4) %>%
  layer_dense(units = 1, activation = "sigmoid")

## Loaded Tensorflow version 2.9.0

nn_1layer %>% compile(loss = "binary_crossentropy",
  optimizer = optimizer_rmsprop(),
  metrics = list("accuracy"))

outcome_train <- as.numeric(viper_train$Outcome)

```

```
nn_fit <- nn_1layer %>% fit(x = x_scale_train,
y = outcome_train)
```

```
nn_preds <- predict(nn_1layer,x=x_scale_test)
nn_class <- if_else(nn_preds >= 0.5, '0','1')
```

```
library(yardstick)
```

```
## Warning: package 'yardstick' was built under R version 4.1.3
```

```
## For binary classification, the first factor level is assumed to be the event.
## Use the argument 'event_level = "second"' to alter this as needed.
```

```
##
```

```
## Attaching package: 'yardstick'
```

```
## The following object is masked from 'package:readr':
```

```
##
```

```
##      spec
```

```
## The following object is masked from 'package:keras':
```

```
##
```

```
##      get_weights
```

```
prediction_df <- data.frame(
  actual = viper_test$Outcome,
  rf = viper_class,
  nn = factor(nn_class, levels = c('0','1'))
)
```

- Further investigating the prediction accuracy of at least one model (e.g., by creating a confusion matrix or a ROC Curve, computing sensitivity/specificity, etc.)

```
#nnet
```

```
#nn
```

```
(nn_conf<- conf_mat(prediction_df, truth = actual, estimate = nn))
```

```
##           Truth
## Prediction    0    1
##           0 103  15
##           1   17   5
```

```
(rf_conf<- conf_mat(prediction_df, truth = actual, estimate = rf))
```

```
##           Truth
## Prediction    0    1
##           0 113  14
##           1   7   6
```

- Selecting a best model and justifying your choice; in particular, you should explain why your model is better than a very stupid model that predicts no one will die

My best model is accurate 85% of the time, but at least it is still better than predicting no one will die(?)

```
summary(nn_conf, event_level = "second")
```

```
## # A tibble: 13 x 3
##   .metric      .estimator .estimate
##   <chr>        <chr>      <dbl>
## 1 accuracy    binary      0.771
## 2 kap         binary      0.104
## 3 sens        binary      0.25
## 4 spec        binary      0.858
## 5 ppv         binary      0.227
## 6 npv         binary      0.873
## 7 mcc         binary      0.104
## 8 j_index     binary      0.108
## 9 bal_accuracy binary      0.554
## 10 detection_prevalence binary      0.157
## 11 precision   binary      0.227
## 12 recall      binary      0.25
## 13 f_meas      binary      0.238
```

```
table(prediction_df$actual, prediction_df$nn)
```

```
##
##      0    1
## 0 103  17
## 1   15   5
```

```
summary(rf_conf, event_level = 'second')
```

```
## # A tibble: 13 x 3
##   .metric      .estimator .estimate
##   <chr>        <chr>      <dbl>
## 1 accuracy    binary      0.85
## 2 kap         binary      0.283
## 3 sens        binary      0.3
## 4 spec        binary      0.942
## 5 ppv         binary      0.462
## 6 npv         binary      0.890
## 7 mcc         binary      0.291
## 8 j_index     binary      0.242
## 9 bal_accuracy binary      0.621
## 10 detection_prevalence binary      0.0929
## 11 precision   binary      0.462
## 12 recall      binary      0.3
## 13 f_meas      binary      0.364
```